

## 13.4 Locks

---

- Transactions must be scheduled so that their effect on shared objects is serially equivalent
- A server can achieve serial equivalence by serialising access to objects, e.g. by the use of locks
- for serial equivalence,
  - (a) all access by a transaction to a particular object must be serialized with respect to another transaction's access.
  - (b) all pairs of conflicting operations of two transactions should be executed in the same order.
- to ensure (b), a transaction is not allowed any new locks after it has released a lock
  - **Two-phase locking** - has a 'growing' and a 'shrinking' phase

## Transactions T and U with exclusive locks (Fig. 13.14, same as 13.7)

Transaction T:		Transaction U:	
$balance = b.getBalance()$ $b.setBalance(bal*1.1)$ $a.withdraw(bal/10)$		$balance = b.getBalance()$ $b.setBalance(bal*1.1)$ $c.withdraw(bal/10)$	
Operations	Locks	Operations	Locks
$openTransaction$		$openTransaction$	
$bal = b.getBalance()$	lock B	when U is about to use B, it is still locked for T and U waits	
$b.setBalance(bal*1.1)$			
$a.withdraw(bal/10)$	lock A	$bal = b.getBalance()$	waits for T's lock on B
$closeTransaction$	unlock A, B	...	
when T is about to use B, it is locked for T		lock B	U can now continue
		$b.setBalance(bal*1.1)$	
when T commits, it unlocks B		$c.withdraw(bal/10)$	lock C
		$closeTransaction$	unlock B, C

- initially the balances of A, B and C unlocked
- the use of the lock on B effectively serialises access to B

# Strict two-phase locking

---

- strict executions prevent **dirty reads** and **premature writes** (if transactions abort).
  - a transaction that **reads** or **writes** an object must be delayed until other transactions that wrote the same object have committed or aborted.
  - to enforce this, any locks applied during the progress of a transaction are held until the transaction commits or aborts.
  - this is called ***strict two-phase locking***
  - For recovery purposes, locks are held until updated objects have been written to permanent storage
- granularity - apply locks to small things e.g. bank balances
  - there are no assumptions as to granularity in the schemes we present

## Read-write conflict rules

---

- concurrency control protocols are designed to deal with *conflicts* between operations in different transactions on the same object
- we describe the protocols in terms of *read* and *write* operations, which we assume are atomic
- read operations of different transactions do not conflict
- therefore exclusive locks reduce concurrency more than necessary
- The 'many reader/ single writer' scheme allows several transactions to read an object or a single transaction to write it (but not both)
- It uses read locks and write locks
  - read locks are sometimes called shared locks

# Lock compatibility

- The operation conflict rules tell us that:
  1. If a transaction  $T$  has already performed a *read* operation on a particular object, then a concurrent transaction  $U$  must not *write* that object until  $T$  commits or aborts.
  2. If a transaction  $T$  has already performed a *write* operation on a particular object, then a concurrent transaction  $U$  must not *read* or *write* that object until  $T$  commits or aborts.

to enforce 1, a request for a write lock is delayed by the presence of a read lock belonging to another transaction

to enforce 2, a request for a read lock or write lock is delayed by the presence of a write lock belonging to another transaction

For one object		Lock requested	
		read	write
Lock already set	none	OK	OK
	read	OK	wait
	write	wait	wait

Figure 13.15

## Lock promotion

---

- Lost updates – two transactions read an object and then use it to calculate a new value.
- Lost updates are prevented by making later transactions delay their reads until the earlier ones have completed.
- each transaction sets a read lock when it reads and then promotes it to a write lock when it writes the same object
- when another transaction requires a read lock it will be delayed until any current transaction has completed
- Lock promotion: the conversion of a lock to a stronger lock – that is, a lock that is more exclusive.
  - demotion of locks (making them weaker) is not allowed
  - The result will be more permissive than the previous one and may allow executions by other transactions that are inconsistent with serial equivalence

# Use of locks in strict two-phase locking

1. When an operation accesses an object within a transaction:
  - (a) If the object is not already locked, it is locked and the operation proceeds.
  - (b) If the object has a conflicting lock set by another transaction, the transaction must wait until it is unlocked.
  - (c) If the object has a non-conflicting lock set by another transaction, the lock is shared and the operation proceeds.
  - (d) If the object has already been locked in the same transaction, the lock will be promoted if necessary and the operation proceeds. (Where promotion is prevented by a conflicting lock, rule (b) is used.)
2. When a transaction is committed or aborted, the server unlocks all objects it locked for the transaction.

Figure 13.16

- The server applies locks when the read/write operations are about to be executed
- the server releases a transaction's locks when it commits or aborts

## Lock implementation

---

- The granting of locks will be implemented by a separate object in the server that we call the *lock manager*.
- the lock manager holds a set of locks, for example in a hash table
- each lock is an instance of the class *Lock* (Fig 13.17) and is associated with a particular object.
  - its variables refer to the object, the holder(s) of the lock and its type
- the lock manager code uses *wait* (when an object is locked) and *notify* when the lock is released
- the lock manager provides *setLock* and *unLock* operations for use by the server



## Figure 13.17 lock class

```
public class Lock {  
    private Object object;           // the object being protected by the lock  
    private Vector holders;          // the TIDs of current holders  
    private LockType lockType;       // the current type  
    public synchronized void acquire(TransID trans, LockType aLockType) {  
        while(/*another transaction holds the lock in conflicting mode*/) {  
            try {  
                wait();  
            } catch ( InterruptedException e) { /* ... */ }  
        }  
        if(holders.isEmpty()) { // no TIDs hold lock  
            holders.addElement(trans);  
            lockType = aLockType;  
        } else if(/*another transaction holds the lock, share it*/ ) {  
            if(/* this transaction not a holder*/) holders.addElement(trans);  
        } else if(/* this transaction is a holder but needs a more exclusive lock*/) {  
            lockType.promote();  
        }  
    }  
}
```

Continues on next slide

## Figure 13.17 continued

---

```
public synchronized void release(TransID trans ){  
    holders.removeElement(trans);           // remove this holder  
    // set locktype to none  
    notifyAll();  
    }  
}
```

## Figure 13.18 *LockManager* class

```
public class LockManager {
    private Hashtable theLocks;

    public void setLock(Object object, TransID trans, LockType
lockType){
        Lock foundLock;
        synchronized(this){
            // find the lock associated with object
            // if there isn't one, create it and add to the hashtable
        }
        foundLock.acquire(trans, lockType);
    }

    // synchronize this one because we want to remove all entries
    public synchronized void unLock(TransID trans) {
        Enumeration e = theLocks.elements();
        while(e.hasMoreElements()){
            Lock aLock = (Lock)(e.nextElement());
            if(/* trans is a holder of this lock*/ ) aLock.release(trans);
        }
    }
}
```

## Deadlock with write locks

T accesses  $A \rightarrow B$   
U accesses  $B \rightarrow A$

Transaction <i>T</i>		Transaction <i>U</i>	
Operations	Locks	Operations	Locks
<i>a.deposit(100);</i>	write lock <i>A</i>		
<i>b.withdraw(100)</i>		<i>b.deposit(200)</i>	write lock <i>B</i>
• • •	waits for <i>U</i> 's	<i>a.withdraw(200);</i>	waits for <i>T</i> 's
• • •	lock on <i>B</i>	• • •	lock on <i>A</i>
• • •		• • •	

Figure 13.19

- The *deposit* and *withdraw* methods are atomic. Although they read as well as write, they acquire write locks.
- The lock manager must be designed to deal with deadlocks.

## The wait-for graph for the previous figure

- Definition of deadlock
  - deadlock is a state in which each member of a group of transactions is waiting for some other member to release a lock.
  - a *wait-for graph* can be used to represent the waiting relationships between current transactions
- In a wait-for graph the nodes represent transactions and the edges represent wait-for relationships between transactions






Figure 13.20



## A cycle in a wait-for graph

---



- Suppose a wait-for graph contains a cycle  $T \dots \rightarrow U \rightarrow \dots \rightarrow V \rightarrow T$ 
  - each transaction waits for the next transaction in the cycle
  - all of these transactions are blocked waiting for locks
  - none of the locks can ever be released (the transactions are deadlocked)
  - If one transaction is aborted, then its locks are released and that cycle is broken

## Another wait-for graph





Figure 13.22



- $T$ ,  $U$  and  $V$  share a read lock on  $C$  and
- $W$  holds write lock on  $B$  (which  $V$  is waiting for)
- $T$  and  $W$  then request write locks on  $C$  and deadlock occurs e.g.  $V$  is in two cycles - look on the left

## Deadlock prevention is unrealistic

---

- e.g. lock all of the objects used by a transaction when it starts
  - unnecessarily restricts access to shared resources.
  - it is sometimes impossible to predict at the start of a transaction which objects will be used.
- Deadlock can also be prevented by requesting locks on objects in a predefined order
  - but this can result in premature locking and a reduction in concurrency



# Deadlock detection

---

- by finding cycles in the wait-for graph
  - after detecting a deadlock, a transaction must be selected to be aborted to break the cycle
  - the software for deadlock detection can be part of the lock manager
  - it holds a representation of the wait-for graph so that it can check it for cycles from time to time
  - edges are added to the graph and removed from the graph by the lock manager's *setLock* and *unLock* operations
  - when a cycle is detected, choose a transaction to be aborted and then remove from the graph all the edges belonging to it
  - it is hard to choose a victim - e.g. choose the oldest or the one in the most cycles

## Timeouts on locks

- Lock timeouts can be used to resolve deadlocks
  - each lock is given a limited period in which it is invulnerable.
  - after this time, a lock becomes vulnerable.
  - provided that no other transaction is competing for the locked object, the vulnerable lock is allowed to remain.
  - but if any other transaction is waiting to access the object protected by a vulnerable lock, the lock is broken
    - ♦ (that is, the object is unlocked) and the waiting transaction resumes.
  - The transaction whose lock has been broken is normally aborted
- problems with lock timeouts
  - locks may be broken when there is no deadlock
  - if the system is overloaded, lock timeouts will happen more often and long transactions will be penalised
  - it is hard to select a suitable length for a timeout

## 13.5 Optimistic concurrency control

---

- the scheme is called optimistic because the likelihood of two transactions conflicting is low
- a transaction proceeds without restriction until the *closeTransaction* (no waiting, therefore no deadlock)
- it is then checked to see whether it has come into conflict with other transactions
- when a conflict arises, a transaction is aborted
- each transaction has three phases

# three phases

## Working phase

- the transaction uses a tentative version of the objects it accesses (dirty reads can't occur as we read from a committed version or a copy of it)
- the coordinator records the *readset* and *writeset* of each transaction

## Validation phase

- at *closeTransaction* the coordinator validates the transaction (looks for conflicts)
- if the validation is successful the transaction can commit.
- if it fails, either the current transaction, or one it conflicts with is aborted

## Update phase

- If validated, the changes in its tentative versions are made permanent.
- read-only transactions can commit immediately after passing validation.

## 13.6 Timestamp ordering concurrency control

---

- each operation in a transaction is validated when it is carried out
  - if an operation cannot be validated, the transaction is aborted
  - each transaction is given a unique timestamp when it starts.
    - ♦ The timestamp defines its position in the time sequence of transactions.
  - requests from transactions can be totally ordered by their timestamps.
- basic timestamp ordering rule (based on operation conflicts)
  - A request to write an object is valid only if that object was last read and written by earlier transactions.
  - A request to read an object is valid only if that object was last written by an earlier transaction

# Transaction commits with timestamp ordering

---

- when a coordinator receives a commit request, it will always be able to carry it out because all operations have been checked for consistency with earlier transactions
  - committed versions of an object must be created in timestamp order
  - the server may sometimes need to wait, but the client need not wait
  - to ensure recoverability, the server will save the 'waiting to be committed versions' in permanent storage
- the timestamp ordering algorithm is strict because
  - the read rule delays each read operation until previous transactions that had written the object had committed or aborted
  - writing the committed versions in order ensures that the write operation is delayed until previous transactions that had written the object have committed or aborted

# Comparison of methods for concurrency control

- pessimistic approach (detect conflicts as they arise)
  - timestamp ordering: serialisation order decided statically
  - locking: serialisation order decided dynamically
  - timestamp ordering is better for transactions where reads >> writes,
  - locking is better for transactions where writes >> reads
  - strategy for aborts
    - ♦ timestamp ordering – immediate
    - ♦ locking– waits but can get deadlock
- optimistic methods
  - all transactions proceed, but may need to abort at the end
  - efficient operations when there are few conflicts, but aborts lead to repeating work
- the above methods are not always adequate e.g.
  - in cooperative work there is a need for user notification
  - applications such as cooperative CAD need user involvement in conflict resolution

# Summary

---

- Operation conflicts form a basis for the derivation of concurrency control protocols.
  - protocols ensure serializability and allow for recovery by using strict executions
  - e.g. to avoid cascading aborts
- Three alternative strategies are possible in scheduling an operation in a transaction:
  - (1) to execute it immediately, (2) to delay it, or (3) to abort it
  - strict two-phase locking uses (1) and (2), aborting in the case of deadlock
    - ♦ ordering according to when transactions access common objects
  - timestamp ordering uses all three - no deadlocks
    - ♦ ordering according to the time transactions start.
  - optimistic concurrency control allows transactions to proceed without any form of checking until they are completed.
    - ♦ Validation is carried out. Starvation can occur.